

CSED311 Lab5: Cache

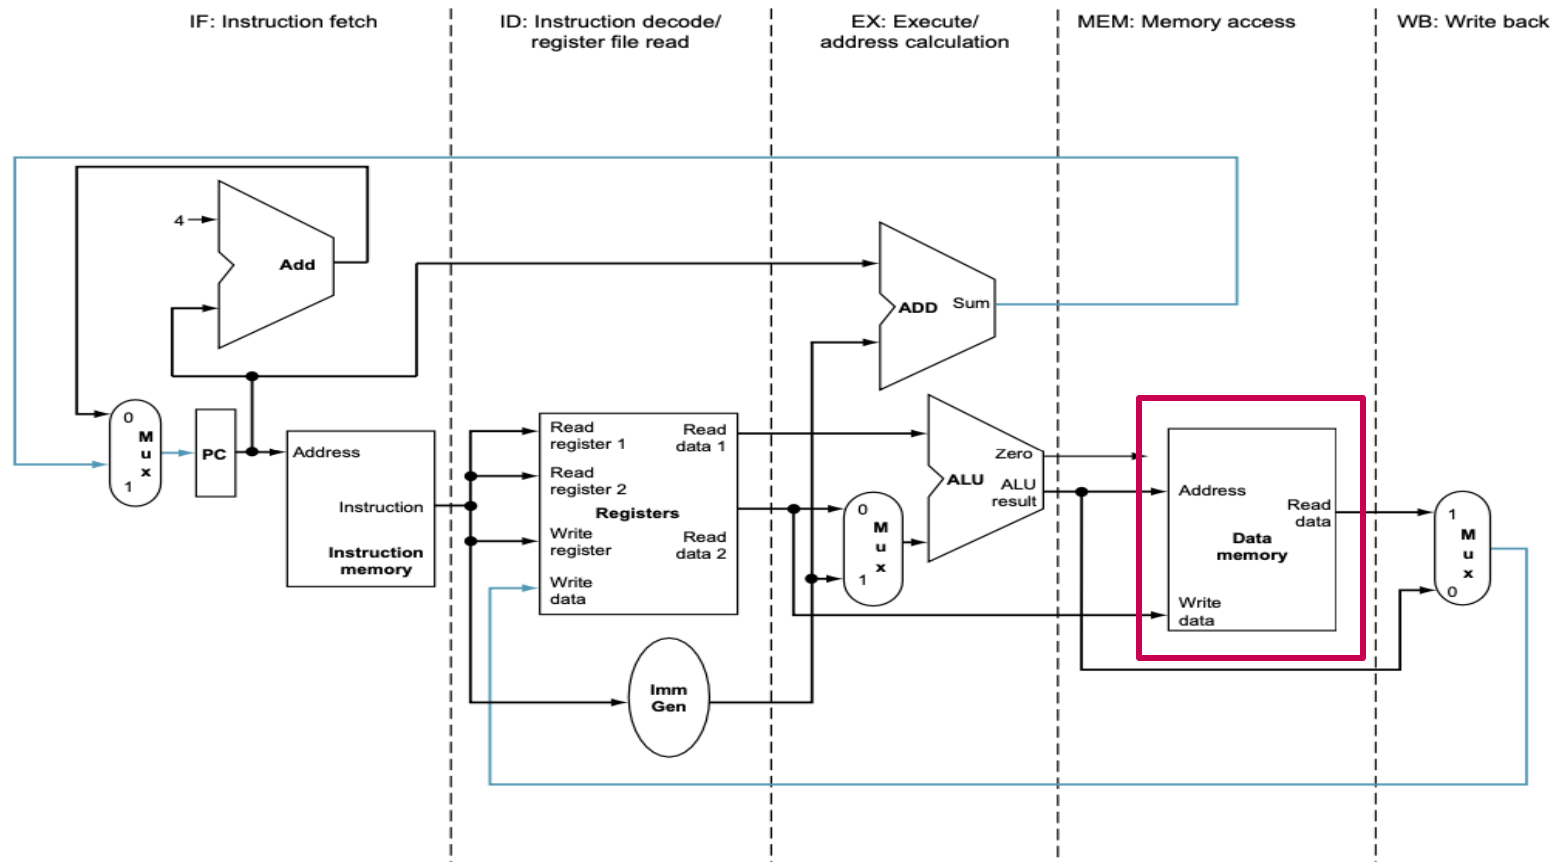
Okkyun Woo

okkyun.w@postech.ac.kr

Contact the TAs at csed311-ta@postech.ac.kr

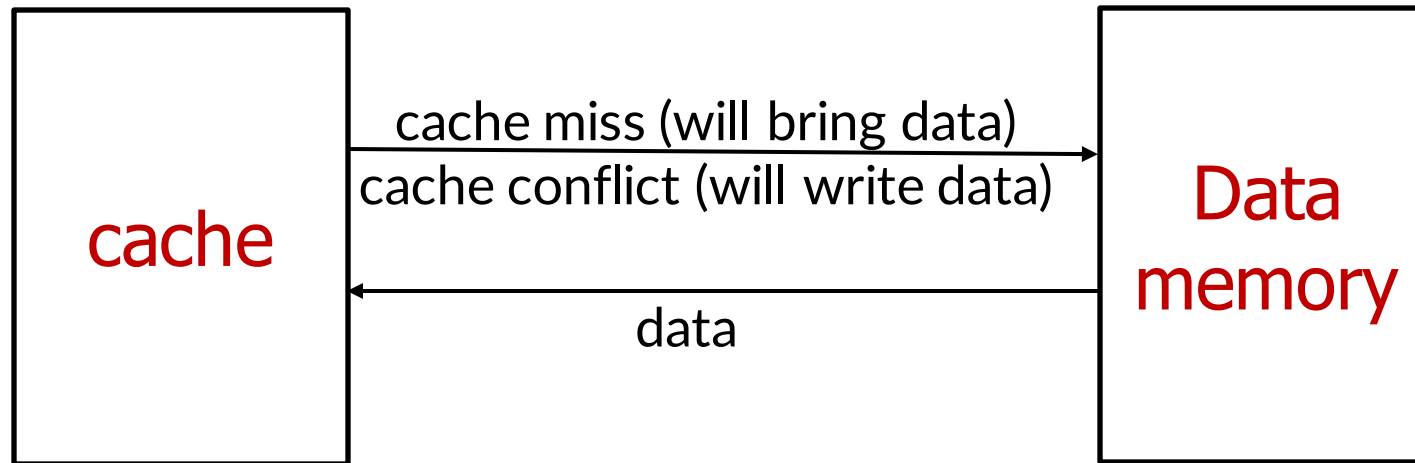
Data cache in Pipelined CPU

- Uses a **blocking data cache** instead of a “magic memory”

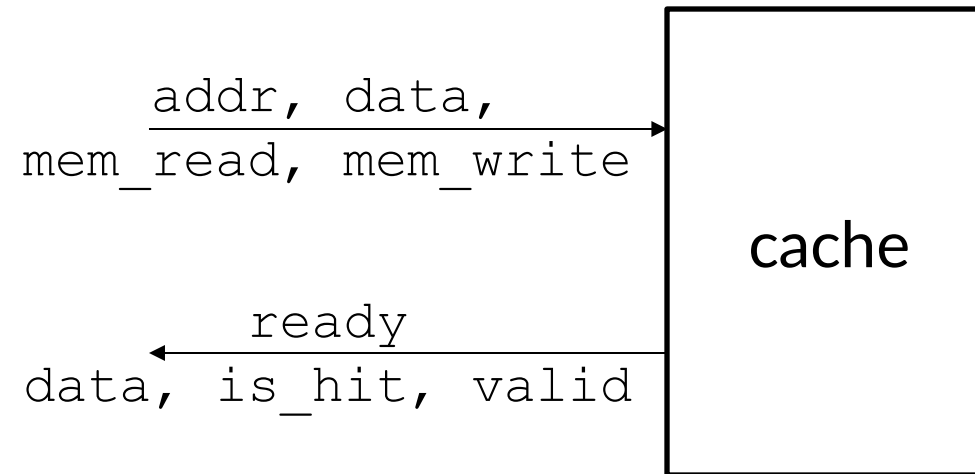


Data flow

- Cache fetches data from the memory

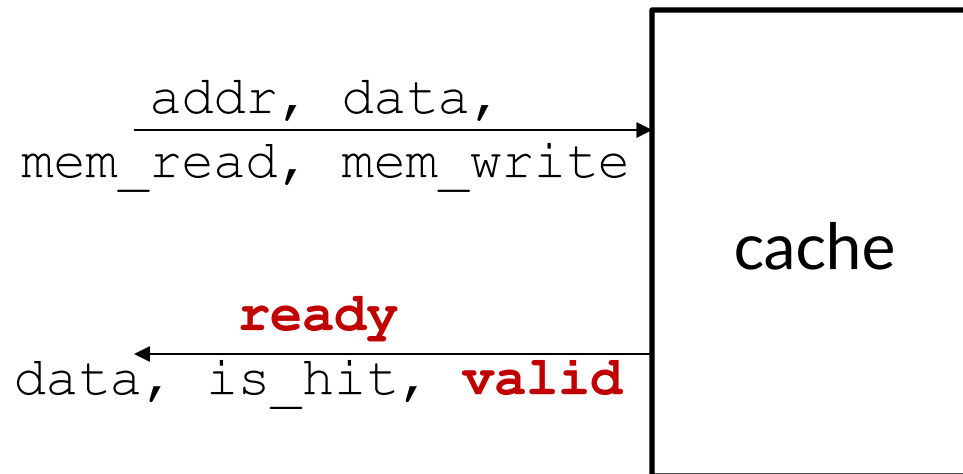


Signals to / from the cache



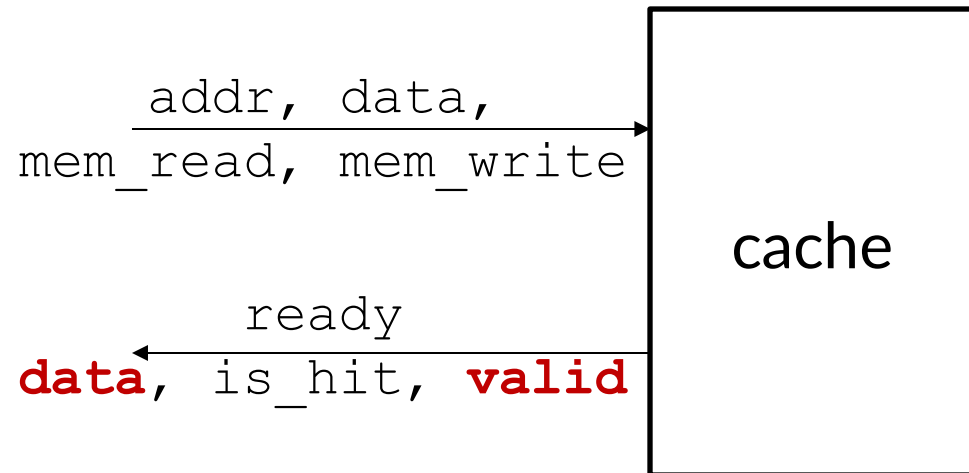
Signals from the data cache

- `ready` represents the status of the cache
 - If `ready` is true => cache is ready to accept request
 - If `ready` is false => cache is busy bringing data from the memory
 - `valid` is false
 - Cache cannot accept a request currently (try later)
 - Need to stall the pipeline



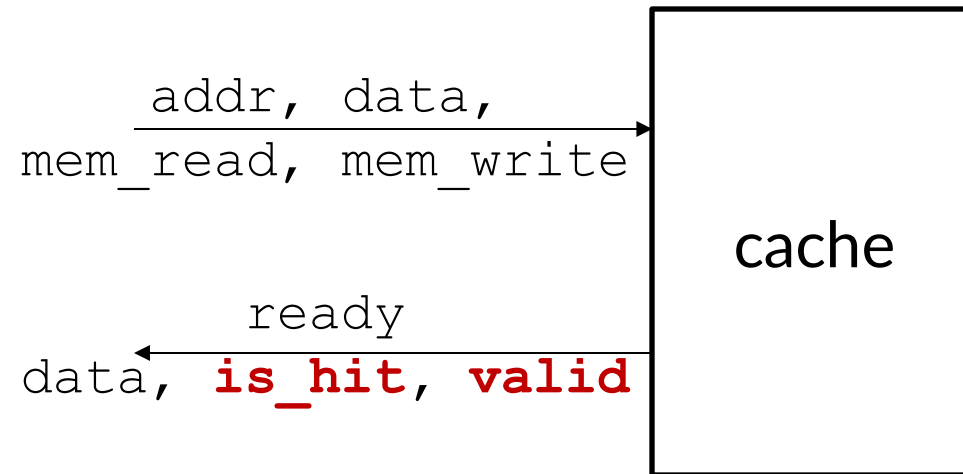
Signals from the data cache

- data represents data
 - Ignore data when valid is false




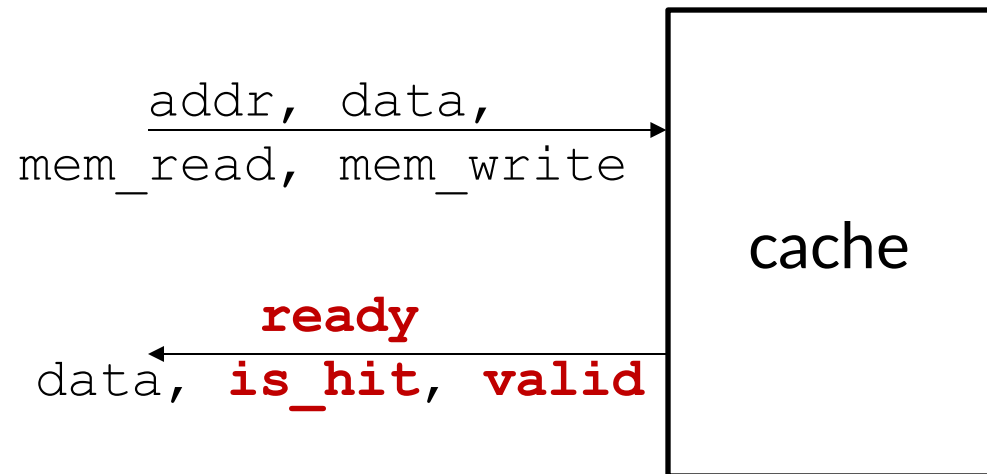
Signals from the data cache

- `is_hit` represents whether a cache hit occurs
 - If a cache miss occurs, stall the pipeline
 - Even if `ready` is true




Signals from the data cache

- Pipeline is not stalled **only when**:
 - **ready, is_hit and valid** are true



Data cache design

- The size of a cache line (block) is 16 bytes



# sets	# ways	Block offset 0	Block offset 1	Block offset 2	Block offset 4
Set 0	Way 0	4B	4B	4B	4B
	Way 1				
	...				

Data cache design


- Asynchronous read:
 - `valid, data, is_hit`
- Synchronous write
 - Writes to the cache line (from both CPU and memory) should be synchronous
- Write-back, write-allocate
 - Read data from the memory if a write miss occurs

Data cache design







- Replacement policy
 - Choose any way except for MRU way
- Structure
 - Choose between direct-mapped or set-associative (extra point) but not fully-associative
 - Size: 256 Bytes (data bank)
 - You are free to define # of ways and sets
- Each cache line should have:
 - Valid bit
 - Dirty bit
 - Bits for replacement
 - ...

Matrix data layout


- Memory layout of the matrix (row-major order)
 - Assume each element of the matrix is 4 B
 - Assume cache line size is 16 B



matrix

address	data
0x00	
0x04	
0x08	
0x0c	
0x10	
0x14	
0x18	
0x1c	

memory




cache

Matrix multiplication


■ Naïve implementation

- Is this cache-friendly? No. Why?


```
// matrix op
for (int m = 0; m < M; ++m) {
    for (int n = 0; n < N; ++n) {
        for (int k = 0; k < K; ++k) {
            c[m][n] += a[m][k] + b[k][n];
        }
    }
}
```



x



=




Matrix multiplication


■ Tiled implementation

- Is this cache-friendly? If yes, why?


```
// matrix op
for (int tile_m = 0; tile_m < M; tile_m += TILE) {
  for (int tile_n = 0; tile_n < N; tile_n += TILE) {
    for (int tile_k = 0; tile_k < K; tile_k += TILE) {
      for (int m = tile_m; m < tile_m + TILE; ++m) {
        for (int n = tile_n; n < tile_n + TILE; ++n) {
          for (int k = tile_k; k < tile_k + TILE; ++k) {
            c[m][n] += a[m][k] + b[k][n];
          }
        }
      }
    }
  }
}
```



x



=




Matrix multiplication


■ Tiled implementation

- Is this cache-friendly? If yes, why?
- Reuse data (in the cache) as much as possible within each tile
 - The tile size is set to the cache line size


```
// matrix op
for (int tile_m = 0; tile_m < M; tile_m += TILE) {
    for (int tile_n = 0; tile_n < N; tile_n += TILE) {
        for (int tile_k = 0; tile_k < K; tile_k += TILE) {
            for (int m = tile_m; m < tile_m + TILE; ++m) {
                for (int n = tile_n; n < tile_n + TILE; ++n) {
                    for (int k = tile_k; k < tile_k + TILE; ++k) {
                        c[m][n] += a[m][k] + b[k][n];
                    }
                }
            }
        }
    }
}
```



x



=



Submission

- Implementation (Deadline: 9:00am, 5/28)
 - Blocking data cache
 - Direct-mapped (no extra credit)
 - N-way associative cache (Full extra credit + 3)
 - You need to follow the rules described in [lab_guide.pdf](#)
- Report (Deadline: 23:59, 5/28)
 - The design of the cache
 - Direct-mapped or associative cache
 - Analyze cache hit ratio
 - If you implement associative cache, compare it with direct-mapped cache
 - Explain your replacement policy
 - Naïve matmul vs optimized matmul
 - Why is the cache hit ratio different between two matmul algorithms?
 - What happens to the cache hit ratio if you change the # of sets and # of ways?

Submission

- Implementation file format
 - .zip file name: Lab5_{team_num}_{student1_id}_{student2_id}.zip
 - Contents of the zip file (only *.v):
 - cpu.v
 - ...
 - Do not include top.v, InstMemory.v, DataMemory, RegisterFile.v, and CLOG2.v
- Report file format
 - Lab5_{team_num}_{student1_id}_{student2_id}.pdf